

## Jes Schultz (she/her)

- Product Manager @ Microsoft
- Based in Louisville, KY
- Runner
- Foodie







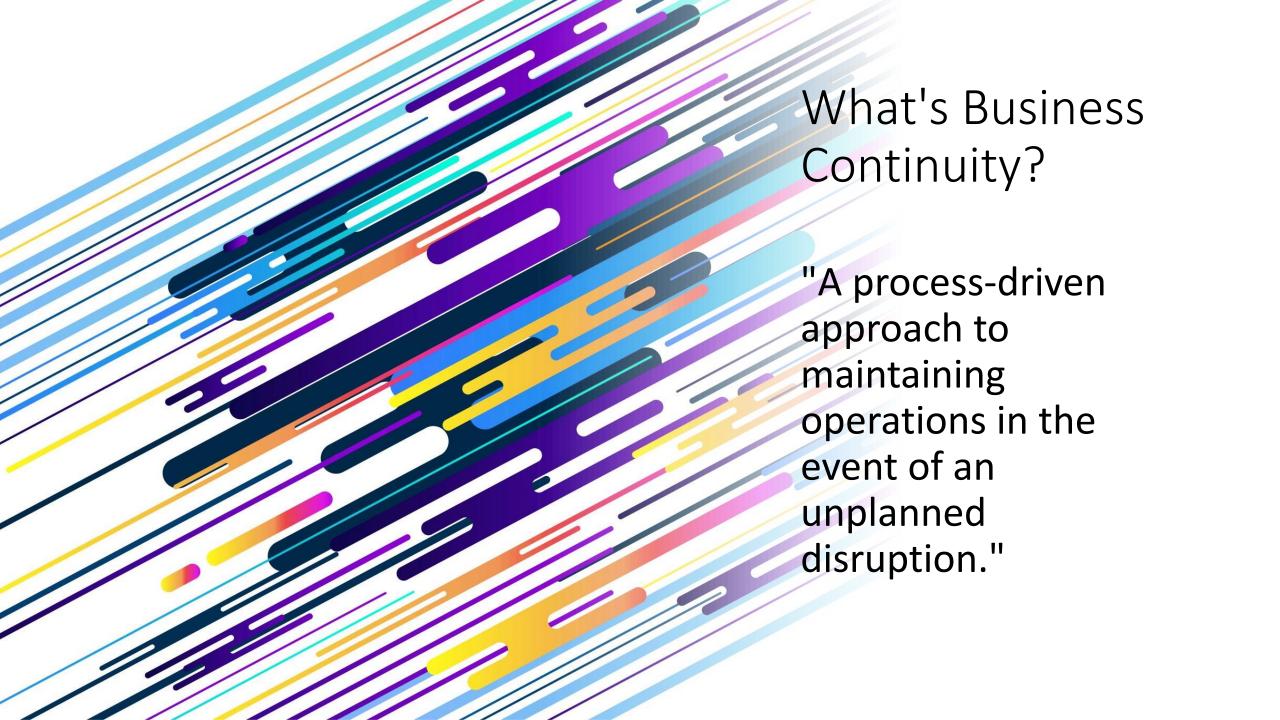






# Data is the most important asset for most organizations.

- If they lose access to their data, or the data itself, they lose business.
- One way to mitigate this threat is through a business continuity plan



# Core concepts

### Backups

 A database backup is a transactionally-consistent record of the database at a point in time

### Restores

 A database restore takes a backup and recreates the database as of a point in time

### HA (High Availability)

Recovering data within the same data center or region

### DR (Disaster Recovery)

Recovering data in a secondary data center or region

## In the on-premises SQL Server world...

- DBAs set up backup jobs
- Hopefully they test restores
- They may set up Windows Server Failover Clustering (WSFC) with SQL Server Failover Cluster Instances (FCI) (shared storage), or Always On Availability Groups (AGs) (separate storage) for HA.
- They may set up Availability Groups or Log Shipping for DR.



### That's a lot of work

...especially when a company has hundreds of SQL Server instances.

## Azure makes it easier!

Azure SQL Database does heavy lifting for the DBA

## Backups





## Taken regularly

Full backup every week

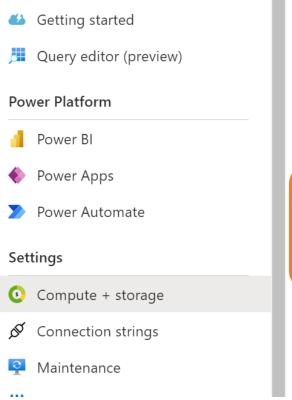
Differential backups every 12 or 24 hours

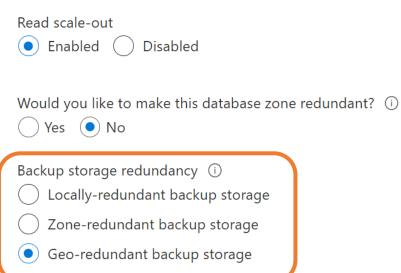
Transaction log backups approximately every 10 minutes - based on compute size and database activity

\*Hyperscale uses storage snapshots

### Redundant

By default, backups are stored in blob storage in the database region and replicated to a paired region





A Selected value for backup storage redundancy is Geo-redundant backup storage. Database backups will be geo-replicated which might impact your data residency requirements. Learn more &

## Restores





Easy button - no more figuring out which series of backups to restore

Go to portal, PowerShell, or CLI and specify the name of the database and the day and time to restore to

### Restore scenarios



RESTORE
EXISTING
DATABASE TO
POINT IN TIME
- NEW
DATABASE,
SAME SERVER,
NEW NAME.



RESTORE
DELETE
DATABASE TO
POINT IN TIME
- ONLY TO
SAME SERVER.



RESTORE DELETED **DATABASE TO ANOTHER** GEOGRAPHIC **REGION - ONLY AVAILABLE FOR DATABASES CONFIGURED** WITH GEO-**REDUNDANT** STODAGE



RESTORE
DATABASE
FROM
SPECIFIC
LONG-TERM
BACKUP

## Demo

Point-in-time restore

# High Availability

Built in – customers don't get to customize in Azure

### ap1 (ap1/ap1) | Compute + storage

SOL database



<<





- Activity log
- Tags
- Diagnose and solve problems
- Getting started
- Query editor (preview)

#### **Power Platform**

- Power BI
- Power Apps
- Power Automate

### Settings

- Compute + storage

### Service and compute tier

Select from the available tiers based on the needs of your workload. The vCore model provides a wide range of configuration controls and offers Hyperscale and Serverless to automatically scale your database based on your workload needs. Alternately, the DTU model provides set price/performance packages to choose from for easy configuration. Learn more

Service tier

General Purpose (Scalable compute and storage options)

vCore-based purchasing model

Compute tier

General Purpose (Scalable compute and storage options)

Hyperscale (On-demand scalable storage)

Business Critical (High transaction rate and high resiliency)

DTU-based purchasing model

Select the hardware configuration based on confidential computing hardware depends (Standard (For workloads with typical performance requirements)

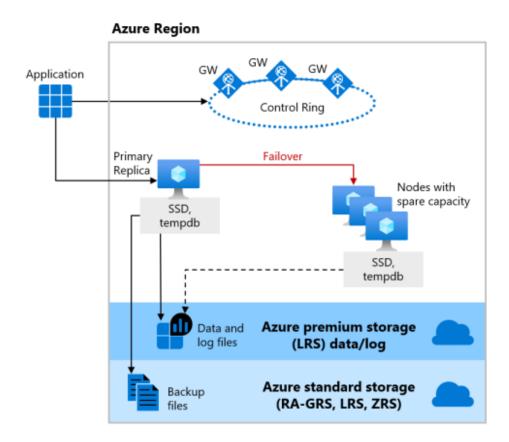
Hardware Configuration

Premium (For IO-intensive workloads)

Change configuration

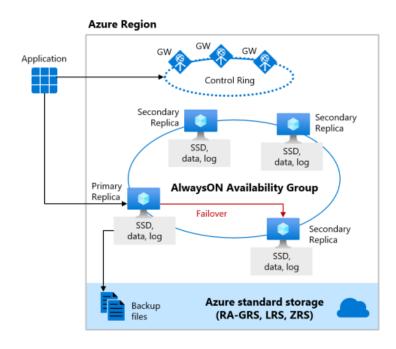
up to 80 vCores, up to 408 GB memory

# General Purpose (Basic, Standard)



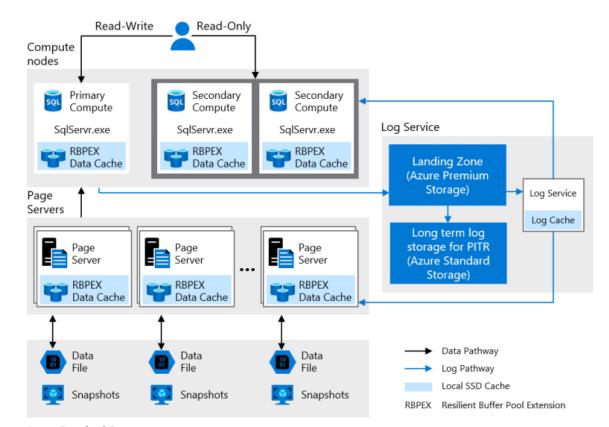
- Separate compute and storage
- The high availability is at the storage level
- During failover, a new stateless compute replica is created and data and log files are attached to it
- Slower failover because the new compute node must start the sqlservr.exe process with a cold cache

# Business Critical (Premium)



- Integrated compute and local SSD storage in a cluster
- At least three nodes exist in the same region
- During failover, primary fails over to a fully-synchronized secondary
- Faster failover because compute is already provisioned, and data is in sync

# Hyperscale



Azure Standard Storage

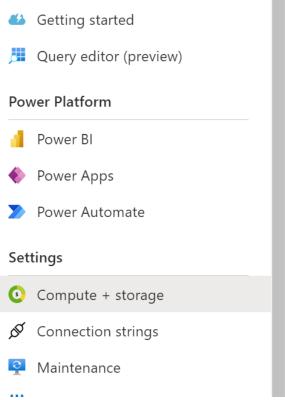
- 4 layers of redundancy
- Compute

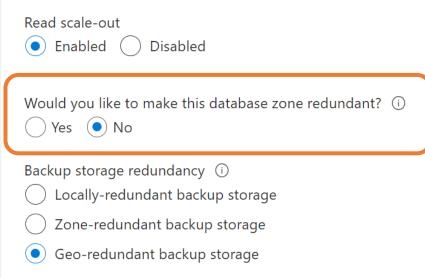
+

- Can have multiple nodes as failover targets
- Page Server storage
  - Every page server has an active-active paired server
- Transaction Log storage
  - Uses Azure Storage for availability and redundancy
- Database file storage
  - Uses Azure Storage for availability and redundancy

## Availability Zones

- By default, all compute nodes and data storage are provisioned in the same datacenter in one region
- Availability Zones (new, still in preview for Hyperscale) allow the compute and storage to be provisioned across multiple data centers in the same region
  - Reduces single point of failure during an outage or disaster
  - Helps companies that need to follow data residency laws achieve HA and some level of DR without going outside of their boundaries





A Selected value for backup storage redundancy is Geo-redundant backup storage. Database backups will be geo-replicated which might impact your data residency requirements. Learn more 

✓

## What else DBAs and devs need to know





The database **can** and **will be** failed over automatically between HA replicas for patching and maintenance.

All apps need to have retry logic built in.

# Disaster Recovery

Customers can configure this!

## Active geo-replication

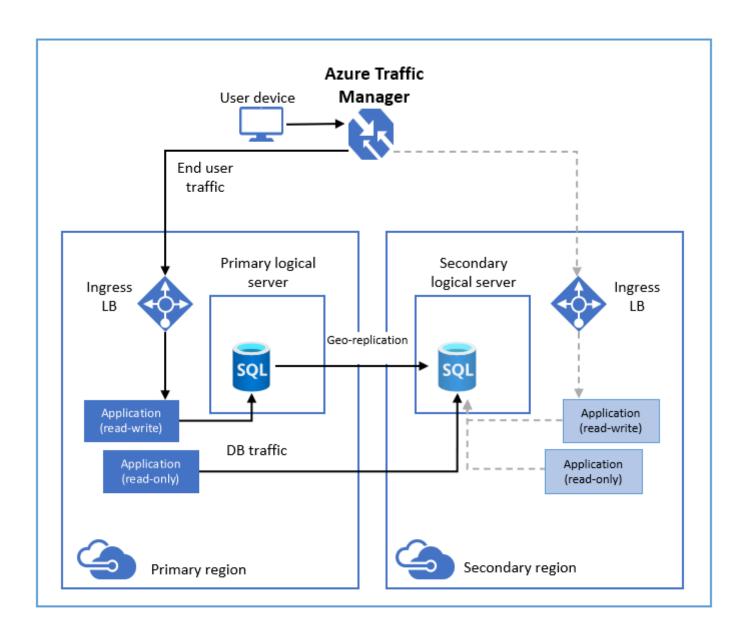
Set up at SQL Database level – one database at a time

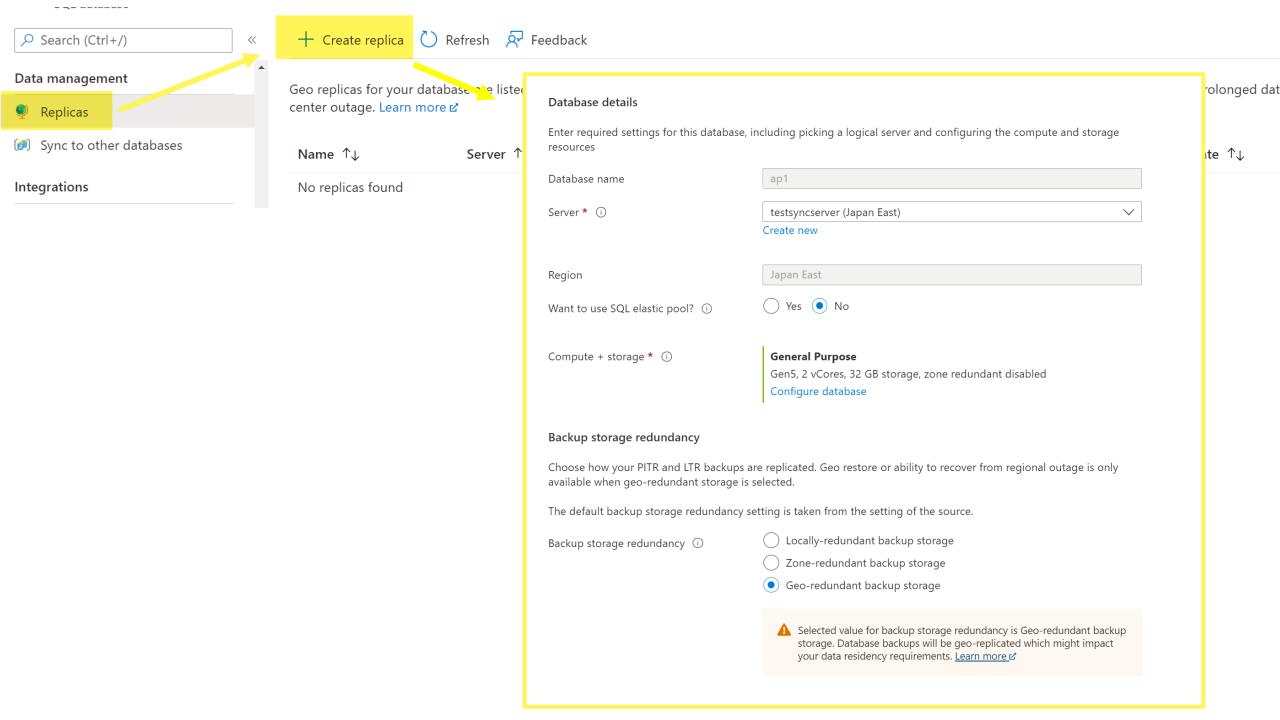
Create up to four readable secondary databases in different region(s)

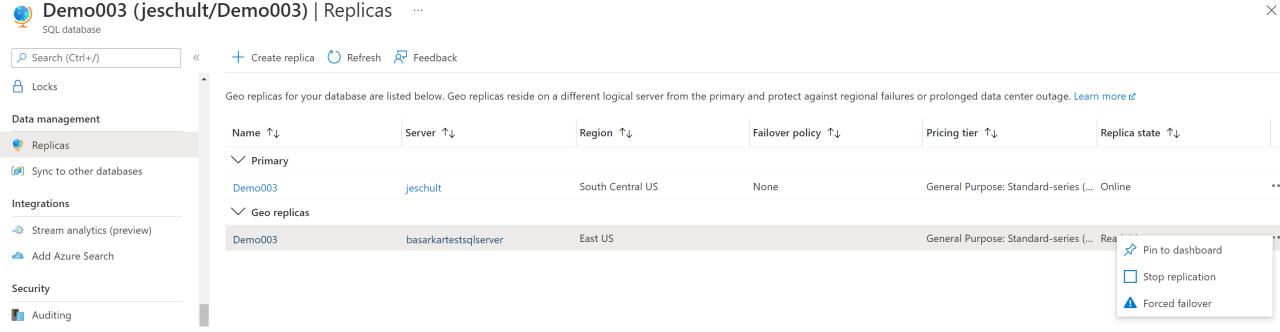
Asynchronous replication from primary to secondaries

Failover: manual

On failover, update application connection string to new server name







## Demo

Active geo-replication

## Why choose active geo-replication?



The app has one database



The app can benefit from one or more readable secondary replicas



The RTO is high enough to allow for manual failover to the replica and time to update the connection string

## Failover groups

ABSTRACTION OF GEO-REPLICATION

SET UP AT SQL SERVER LEVEL!

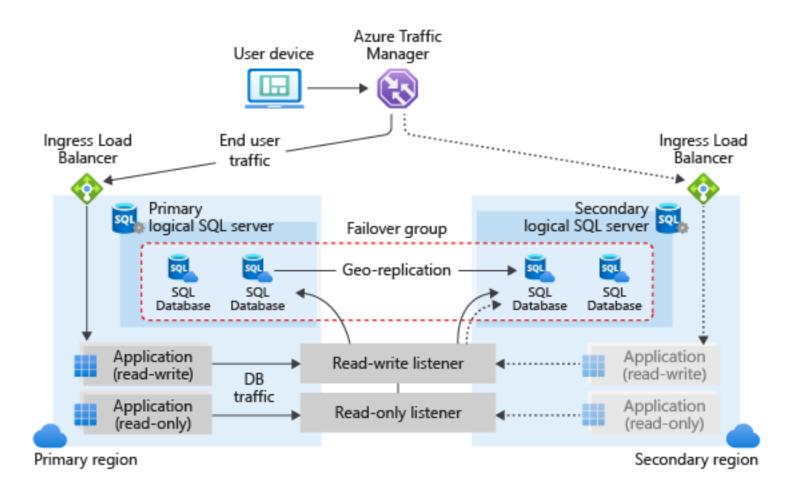
ONE OR MORE DATABASES IN A GROUP

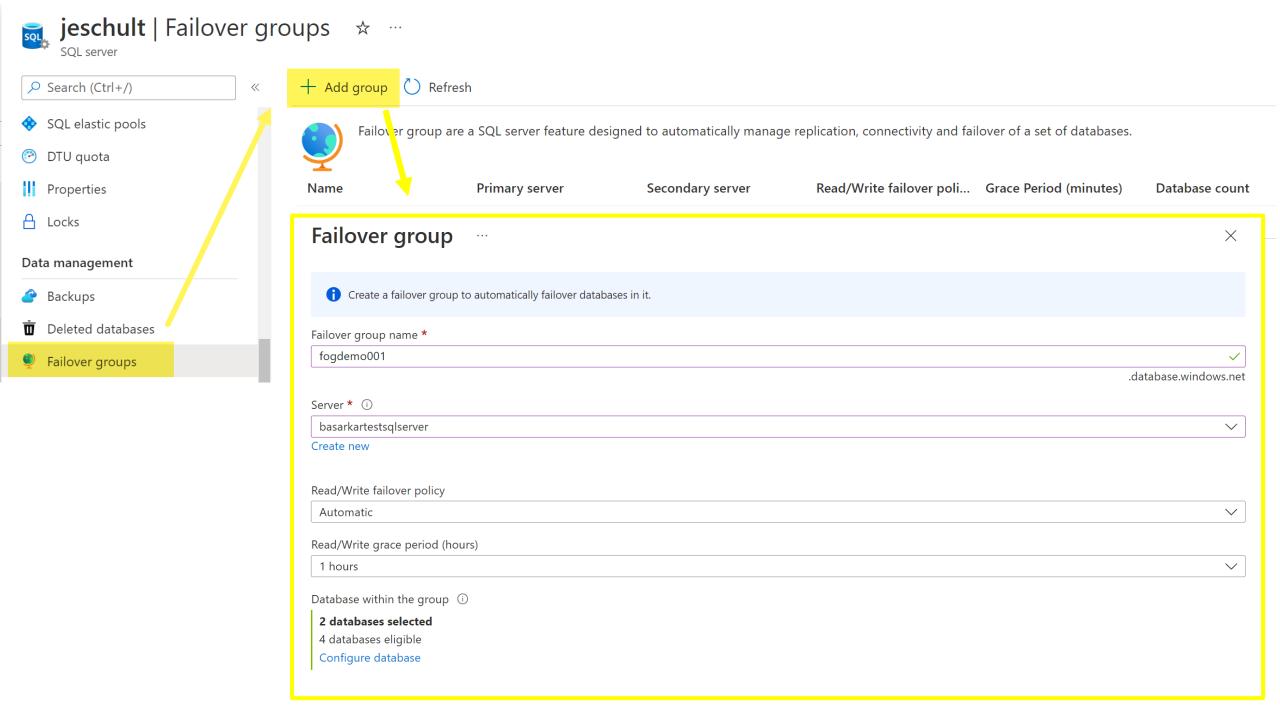
READABLE SECONDARIES BEHIND A LOAD BALANCER WITH A LISTENER

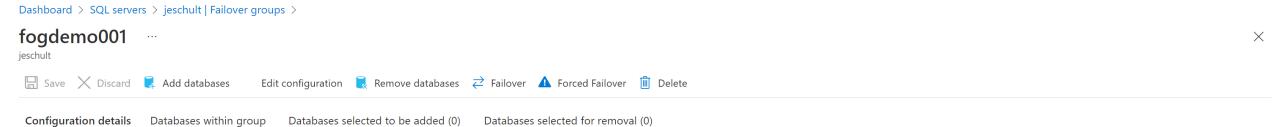
CAN HAVE A READ-WRITE LISTENER OR READ-ONLY LISTENER

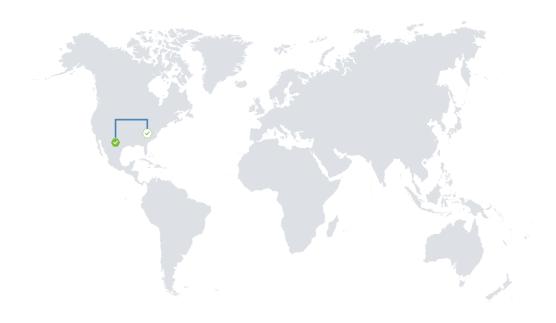
FAILOVER: AUTOMATIC OR MANUAL

ON FAILOVER, IF APPLICATION IS DIRECTED TO LISTENER NAME, NO NEED TO UPDATE CONNECTION STRING









	Server	Role	Read/Write failover policy	Grace period
	jeschult (South Central US)	Primary	Automatic	1 hours
$\langle \checkmark \rangle$	basarkartestsqlserver (East US)	Secondary		

### Read/write listener endpoint

fogdemo001.database.windows.net

### Read-only listener endpoint

# Demo

Failover groups

## Why choose failover groups?



App has multiple databases that need to fail over together



The RTO is low enough that failover must be automatic and the connection string can't be updated manually

# Recap

### Backups

- Azure SQL Database automatically takes backups
- Backups are stored in multiple regions

### Restores

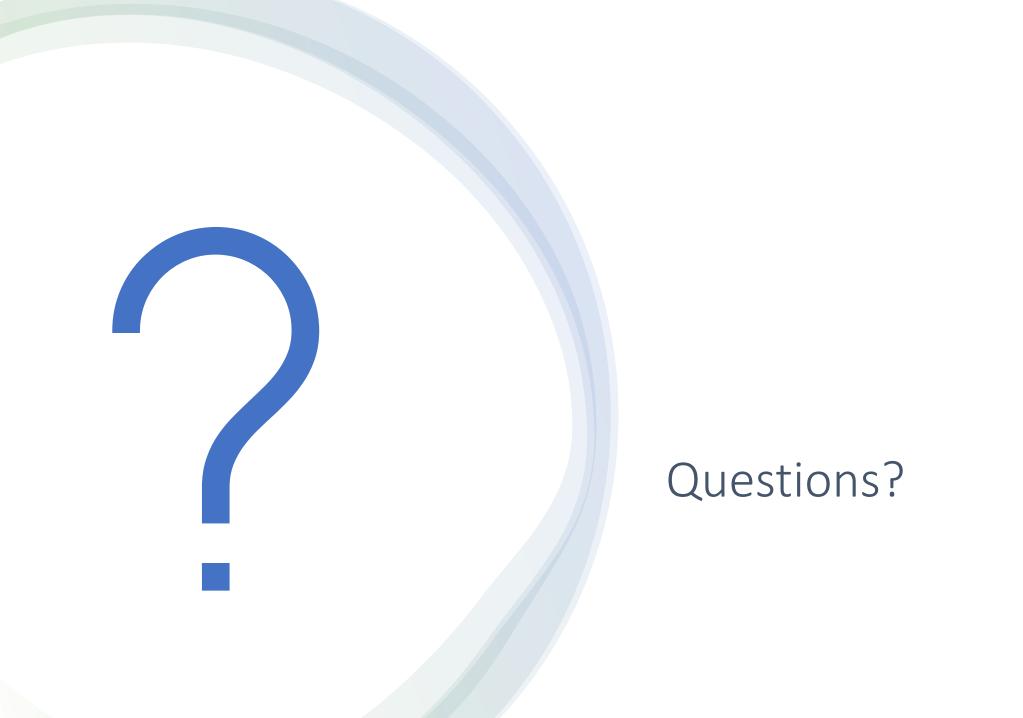
 Azure SQL Database restores are simple – provide the date and time you want to restore to

### HA (High Availability)

• HA is built into every tier of Azure SQL Database

### DR (Disaster Recovery)

- Active geo-replication for a single database
- Failover groups at the SQL server level for multiple databases



### Resources

- Automatic, geo-redundant backups Azure SQL Database | Microsoft
   Learn
- Restore a database from a backup Azure SQL Database | Microsoft Learn
- Active geo-replication Azure SQL Database | Microsoft Learn
- Auto-failover groups overview & best practices Azure SQL Database
   Microsoft Learn
- Disaster recovery drills Azure SQL Database | Microsoft Learn